

Bank Transaction Fraud Detection Using Machine Learning

March 28, 2025

1 Importing Libraries

```
[1]: # Importing necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.ensemble import IsolationForest, RandomForestClassifier
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix
%matplotlib inline

# Setting display options
pd.set_option('display.max_columns', None)
```

2 Loading and Overview of the Dataset

```
[2]: # Loading the dataset
dataset_path = 'bank_transactions_data_2.csv'
df = pd.read_csv(dataset_path)

# Displaying the first five rows of the dataset
print("First five rows of the dataset:")
display(df.head())

# Getting basic information about the dataset
print("\nDataset information:")
print(df.info())

# Checking for missing values
print("\nMissing values in each column:")
print(df.isnull().sum())

# Statistical summary of numerical columns
print("\nStatistical summary of numerical columns:")
```

```

display(df.describe())

# Checking for duplicate entries
print("\nNumber of duplicate rows:")
print(df.duplicated().sum())

```

First five rows of the dataset:

	TransactionID	AccountID	TransactionAmount	TransactionDate	
0	TX000001	AC00128	14.09	2023-04-11 16:29:14	
1	TX000002	AC00455	376.24	2023-06-27 16:44:19	
2	TX000003	AC00019	126.29	2023-07-10 18:16:08	
3	TX000004	AC00070	184.50	2023-05-05 16:32:11	
4	TX000005	AC00411	13.45	2023-10-16 17:51:24	

	TransactionType	Location	DeviceID	IP Address	MerchantID	Channel	
0	Debit	San Diego	D000380	162.198.218.92	M015	ATM	
1	Debit	Houston	D000051	13.149.61.4	M052	ATM	
2	Debit	Mesa	D000235	215.97.143.157	M009	Online	
3	Debit	Raleigh	D000187	200.13.225.150	M002	Online	
4	Credit	Atlanta	D000308	65.164.3.100	M091	Online	

	CustomerAge	CustomerOccupation	TransactionDuration	LoginAttempts	
0	70	Doctor	81	1	
1	68	Doctor	141	1	
2	19	Student	56	1	
3	26	Student	25	1	
4	26	Student	198	1	

	AccountBalance	PreviousTransactionDate	
0	5112.21	2024-11-04 08:08:08	
1	13758.91	2024-11-04 08:09:35	
2	1122.35	2024-11-04 08:07:04	
3	8569.06	2024-11-04 08:09:06	
4	7429.40	2024-11-04 08:06:39	

Dataset information:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 2512 entries, 0 to 2511

Data columns (total 16 columns):

#	Column	Non-Null Count	Dtype
0	TransactionID	2512 non-null	object
1	AccountID	2512 non-null	object
2	TransactionAmount	2512 non-null	float64
3	TransactionDate	2512 non-null	object
4	TransactionType	2512 non-null	object
5	Location	2512 non-null	object

```

6  DeviceID           2512 non-null  object
7  IP Address         2512 non-null  object
8  MerchantID        2512 non-null  object
9  Channel            2512 non-null  object
10 CustomerAge        2512 non-null  int64
11 CustomerOccupation 2512 non-null  object
12 TransactionDuration 2512 non-null  int64
13 LoginAttempts      2512 non-null  int64
14 AccountBalance     2512 non-null  float64
15 PreviousTransactionDate 2512 non-null  object
dtypes: float64(2), int64(3), object(11)
memory usage: 314.1+ KB
None

```

Missing values in each column:

TransactionID	0
AccountID	0
TransactionAmount	0
TransactionDate	0
TransactionType	0
Location	0
DeviceID	0
IP Address	0
MerchantID	0
Channel	0
CustomerAge	0
CustomerOccupation	0
TransactionDuration	0
LoginAttempts	0
AccountBalance	0
PreviousTransactionDate	0

dtype: int64

Statistical summary of numerical columns:

	TransactionAmount	CustomerAge	TransactionDuration	LoginAttempts	\
count	2512.000000	2512.000000	2512.000000	2512.000000	
mean	297.593778	44.673965	119.643312	1.124602	
std	291.946243	17.792198	69.963757	0.602662	
min	0.260000	18.000000	10.000000	1.000000	
25%	81.885000	27.000000	63.000000	1.000000	
50%	211.140000	45.000000	112.500000	1.000000	
75%	414.527500	59.000000	161.000000	1.000000	
max	1919.110000	80.000000	300.000000	5.000000	

	AccountBalance
count	2512.000000
mean	5114.302966

```

std      3900.942499
min     101.250000
25%    1504.370000
50%    4735.510000
75%    7678.820000
max    14977.990000

```

Number of duplicate rows:
0

3 Data Preprocessing

```
[3]: # Converting date columns to datetime objects
df['TransactionDate'] = pd.to_datetime(df['TransactionDate'])
df['PreviousTransactionDate'] = pd.to_datetime(df['PreviousTransactionDate'])

# Extracting new features from date columns
df['TransactionHour'] = df['TransactionDate'].dt.hour
df['TransactionDay'] = df['TransactionDate'].dt.dayofweek # Monday=0, Sunday=6

# Handling missing values (if any)
# Since there are no missing values, we proceed without imputations

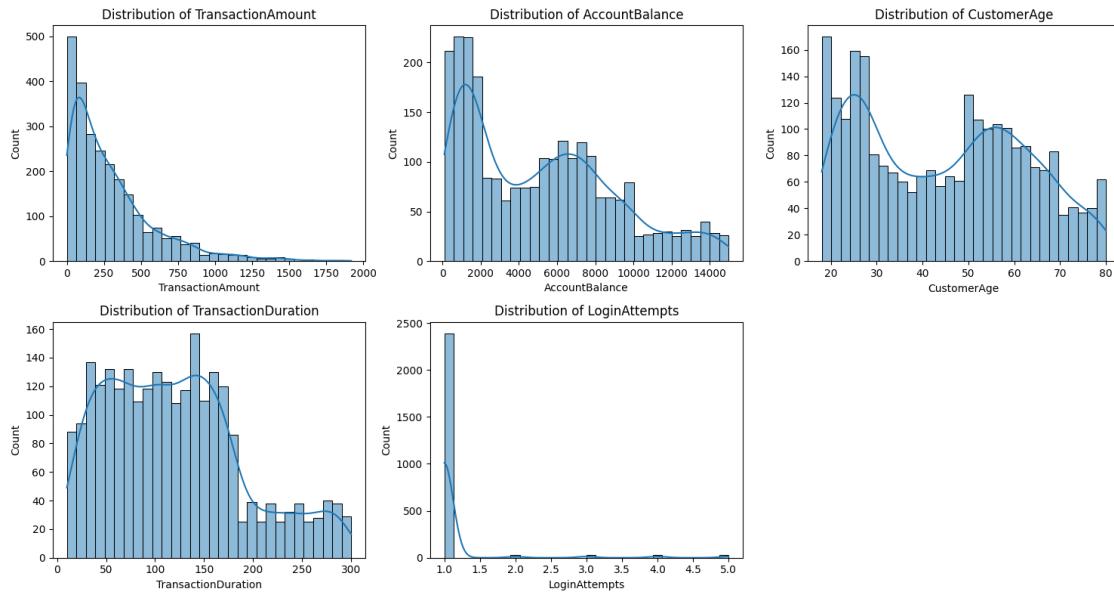
# Checking for negative account balances
if (df['AccountBalance'] < 0).any():
    print("\nTransactions with negative account balance found.")
else:
    print("\nNo transactions with negative account balance.")
```

No transactions with negative account balance.

4 Exploratory Data Analysis (EDA)

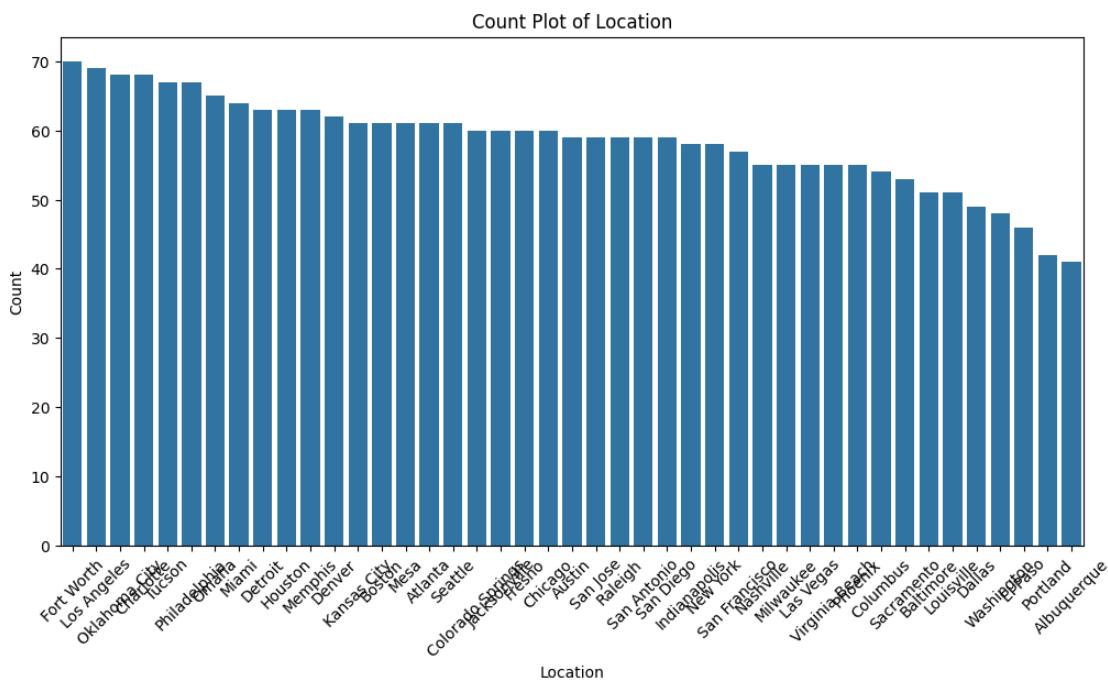
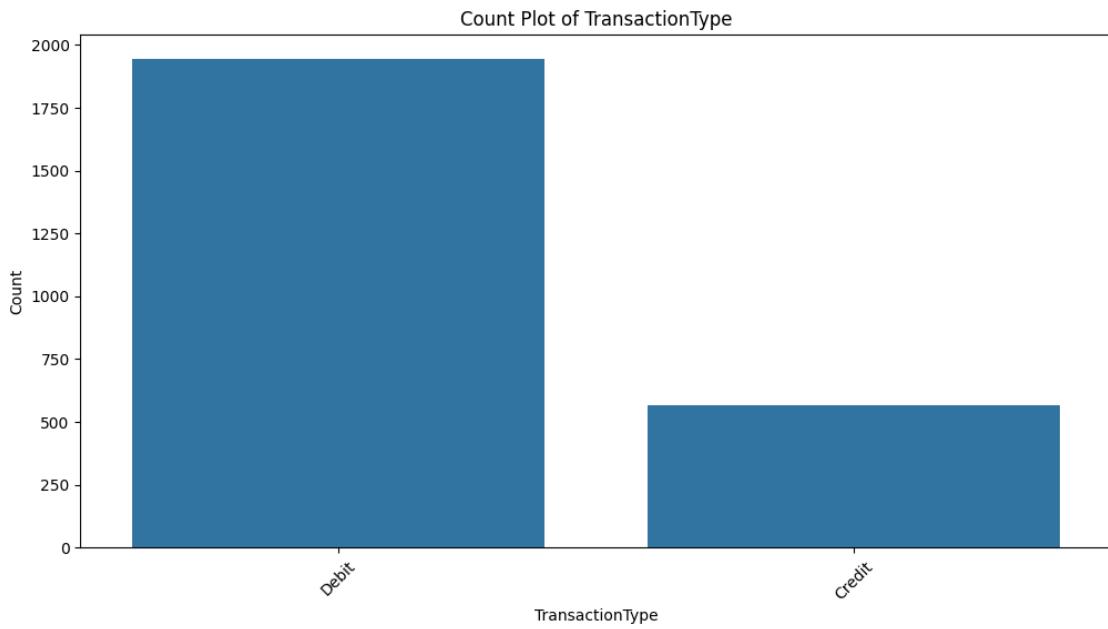
4.1 Plotting distributions of numerical variables

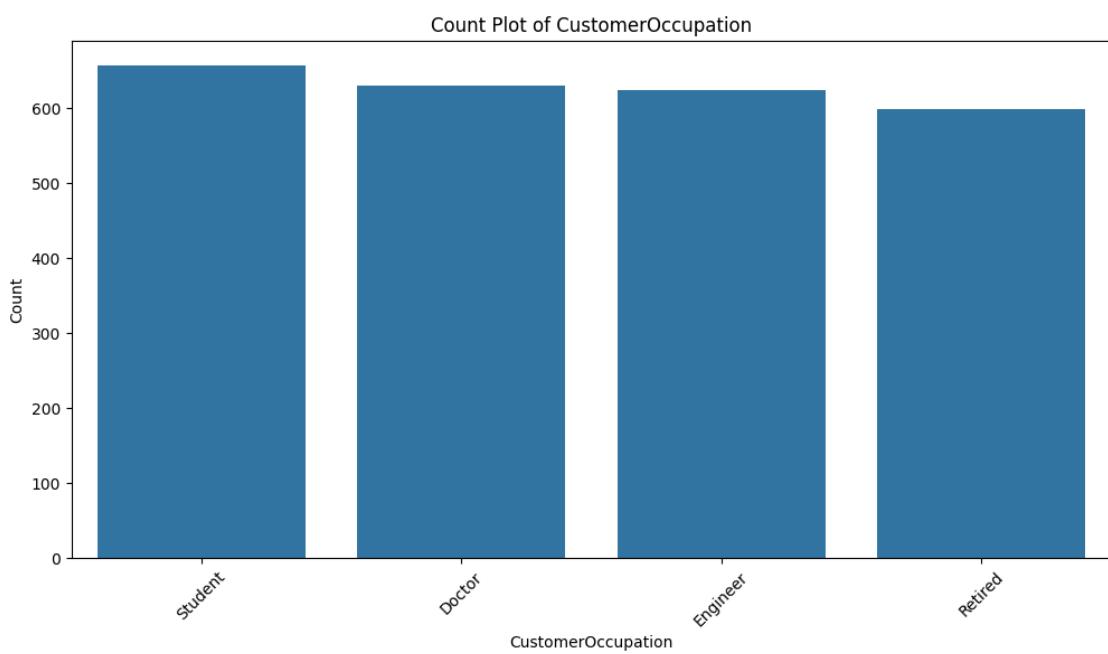
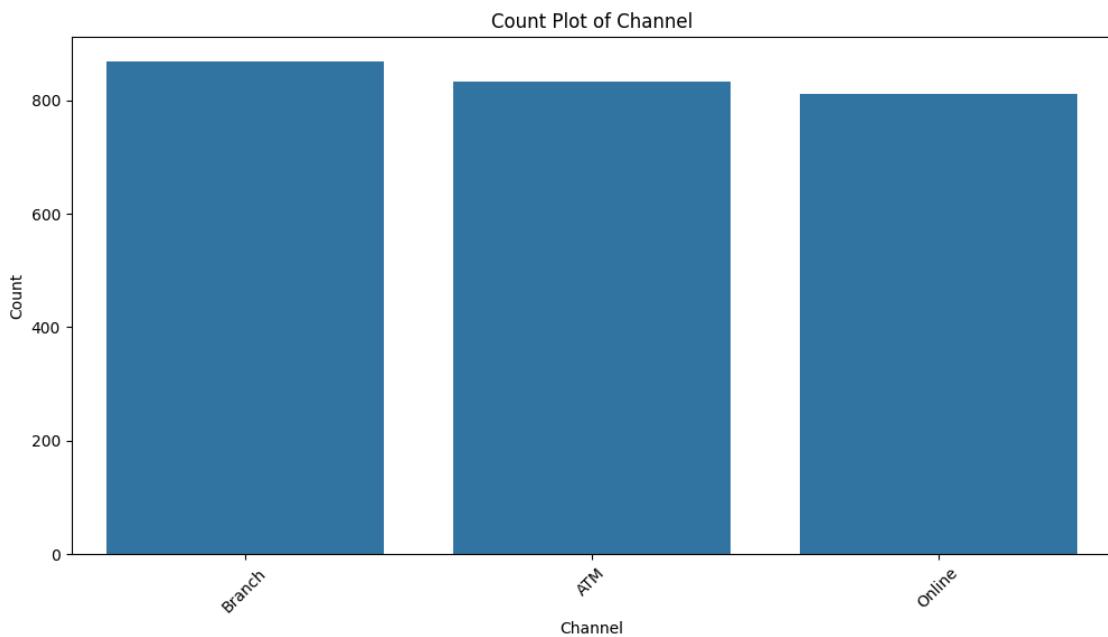
```
[4]: # Univariate Analysis: Plotting distributions of numerical variables
numerical_cols = ['TransactionAmount', 'AccountBalance', 'CustomerAge', 'TransactionDuration', 'LoginAttempts']
plt.figure(figsize=(15, 8))
for i, col in enumerate(numerical_cols):
    plt.subplot(2, 3, i+1)
    sns.histplot(df[col], kde=True, bins=30)
    plt.title(f'Distribution of {col}')
plt.tight_layout()
plt.show()
```



4.2 Categorical Variables Analysis

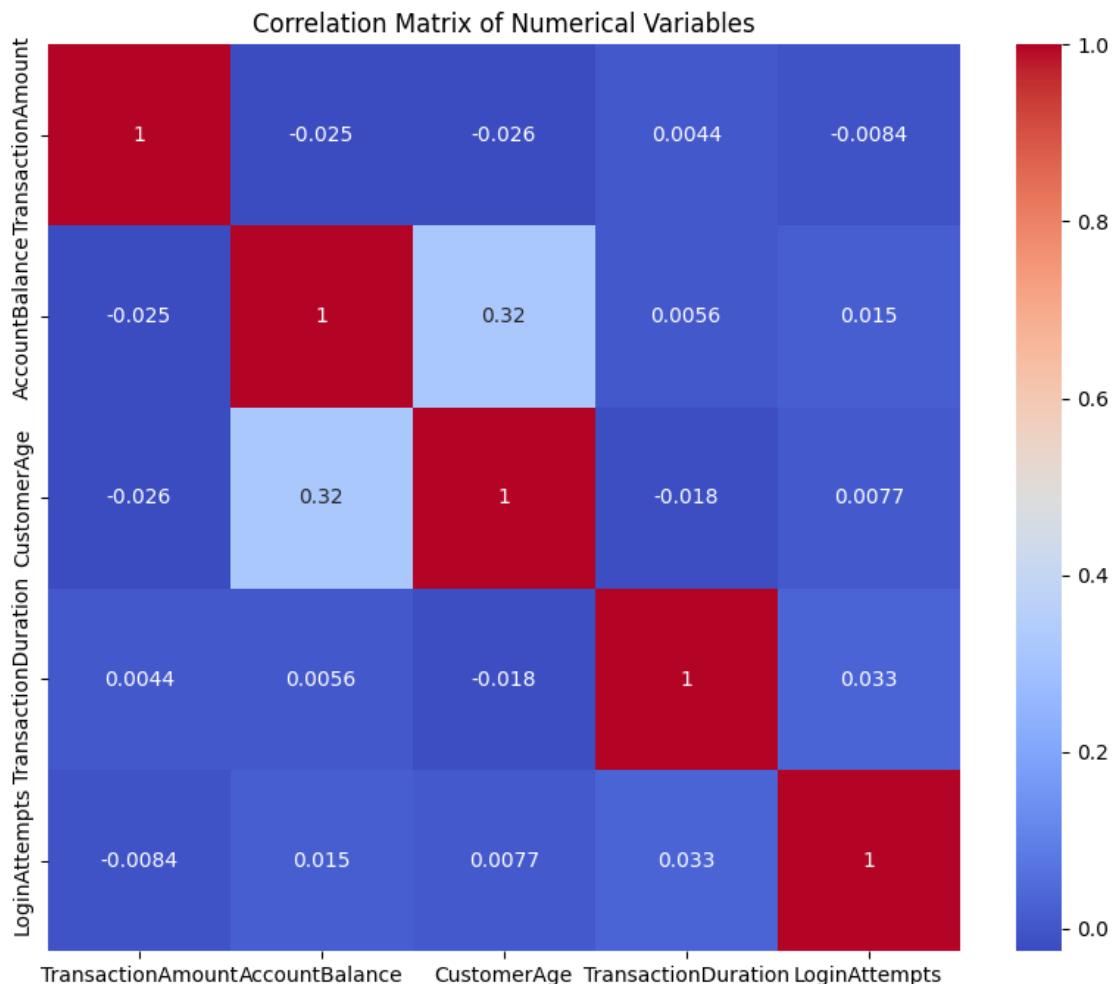
```
[5]: # Categorical Variables Analysis
categorical_cols = ['TransactionType', 'Location', 'Channel', 'CustomerOccupation']
for col in categorical_cols:
    plt.figure(figsize=(12, 6))
    sns.countplot(x=col, data=df, order=df[col].value_counts().index)
    plt.title(f'Count Plot of {col}')
    plt.xlabel(col)
    plt.ylabel('Count')
    plt.xticks(rotation=45)
    plt.show()
```





4.3 Correlation Matrix

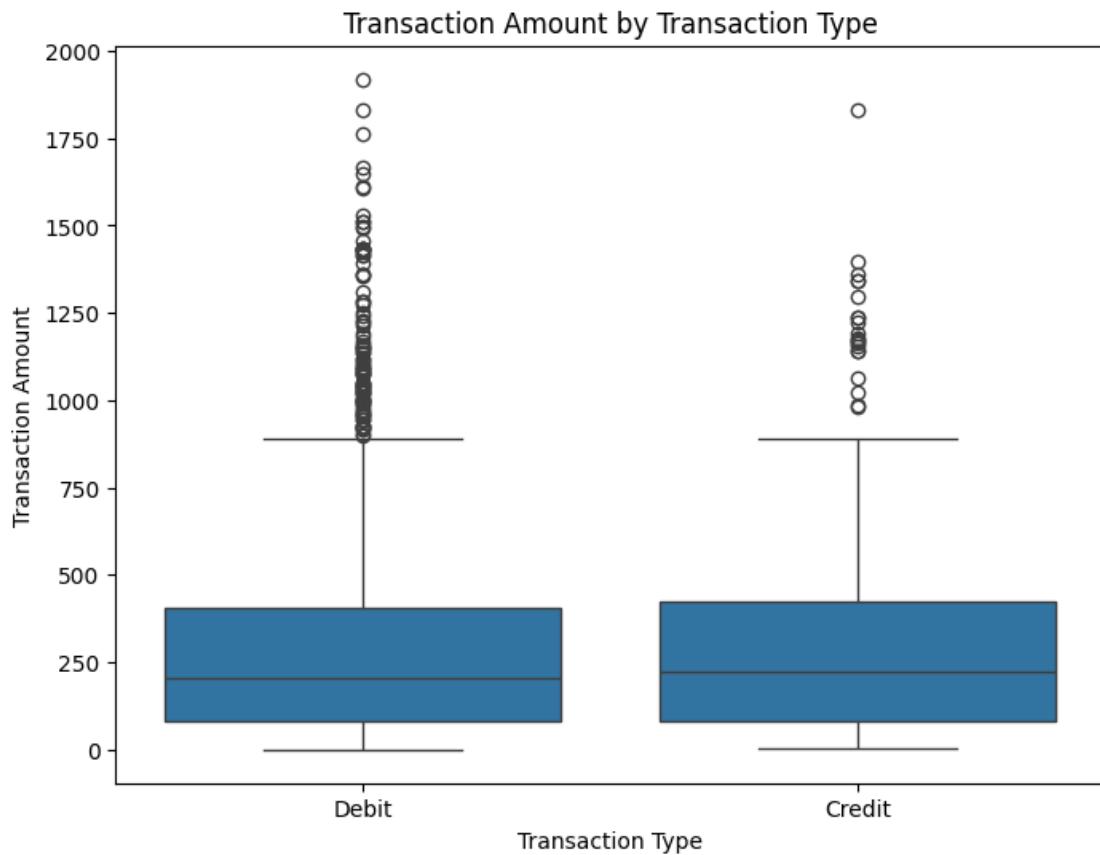
```
[6]: # Correlation Matrix
plt.figure(figsize=(10, 8))
corr_matrix = df[numerical_cols].corr()
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix of Numerical Variables')
plt.show()
```



4.4 Transaction Amount by Transaction Type

```
[7]: # Transaction Amount by Transaction Type
plt.figure(figsize=(8, 6))
sns.boxplot(x='TransactionType', y='TransactionAmount', data=df)
plt.title('Transaction Amount by Transaction Type')
plt.xlabel('Transaction Type')
```

```
plt.ylabel('Transaction Amount')
plt.show()
```



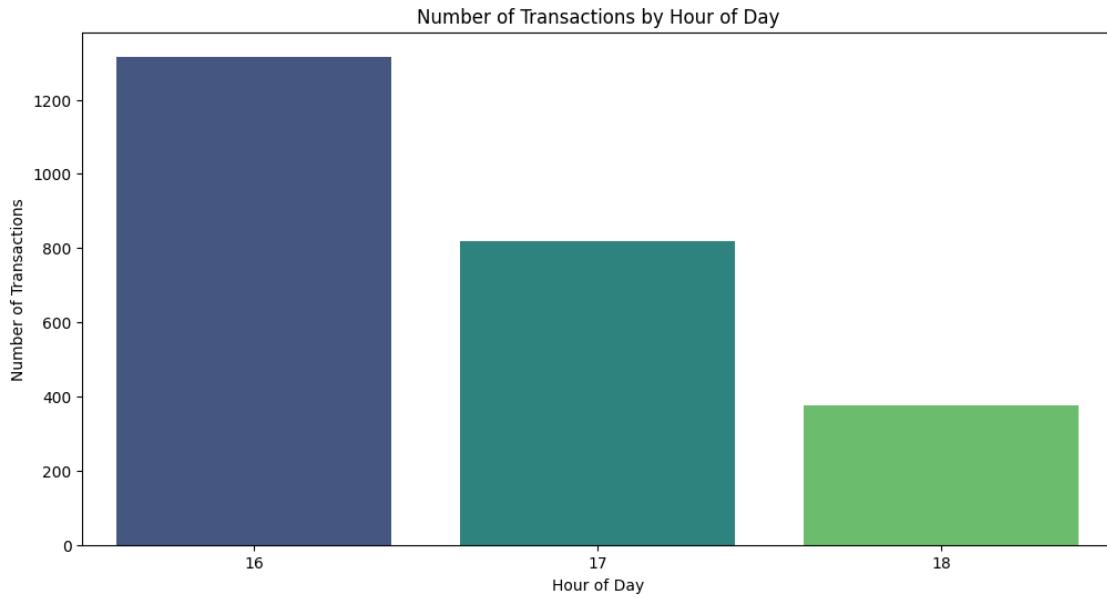
4.5 Number of Transactions by Hour of Day

```
[8]: # Number of Transactions by Hour of Day
plt.figure(figsize=(12, 6))
sns.countplot(x='TransactionHour', data=df, palette='viridis')
plt.title('Number of Transactions by Hour of Day')
plt.xlabel('Hour of Day')
plt.ylabel('Number of Transactions')
plt.show()
```

C:\Users\rakes\AppData\Local\Temp\ipykernel_2976\3355427699.py:3: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(x='TransactionHour', data=df, palette='viridis')
```



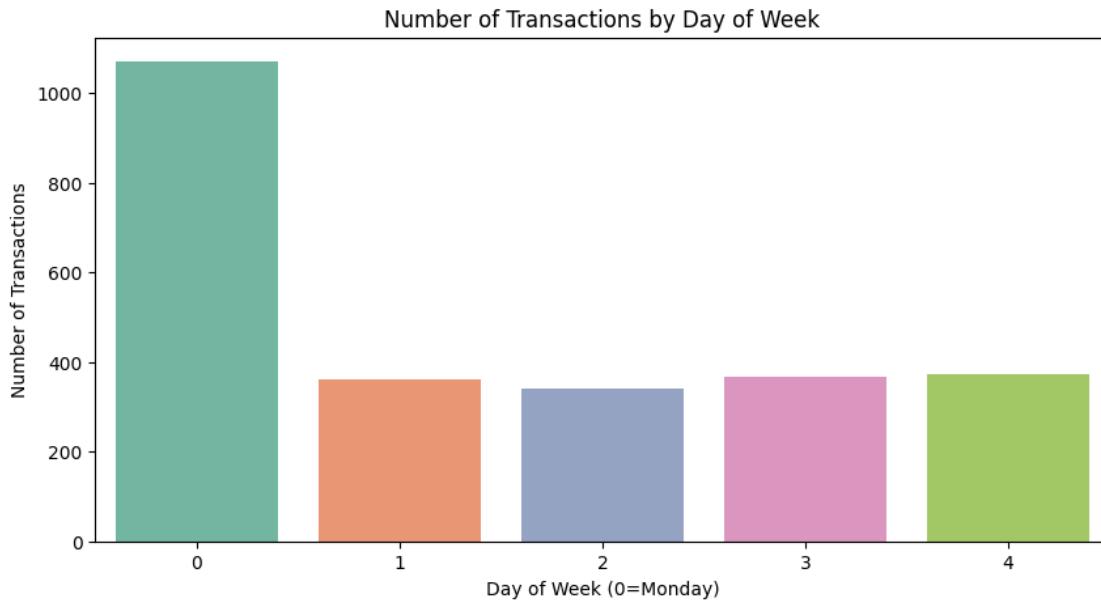
4.6 Number of Transactions by Day of Week

```
[9]: # Number of Transactions by Day of Week
plt.figure(figsize=(10, 5))
sns.countplot(x='TransactionDay', data=df, palette='Set2')
plt.title('Number of Transactions by Day of Week')
plt.xlabel('Day of Week (0=Monday)')
plt.ylabel('Number of Transactions')
plt.show()
```

C:\Users\rakes\AppData\Local\Temp\ipykernel_2976\895084038.py:3: FutureWarning:

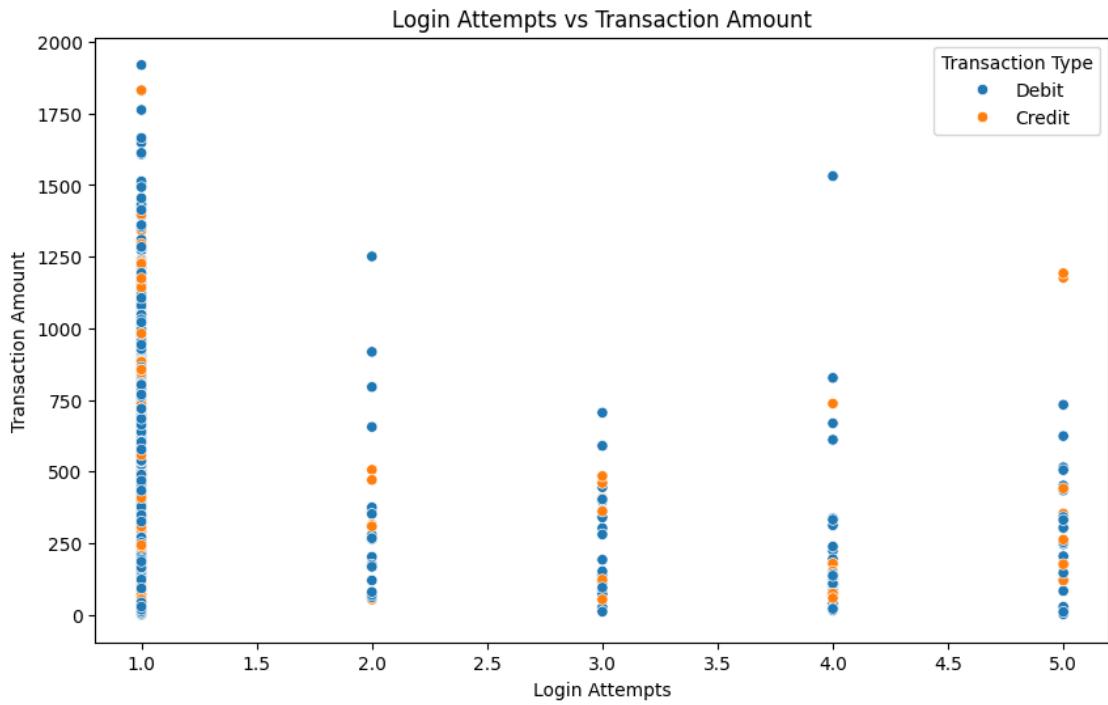
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(x='TransactionDay', data=df, palette='Set2')
```



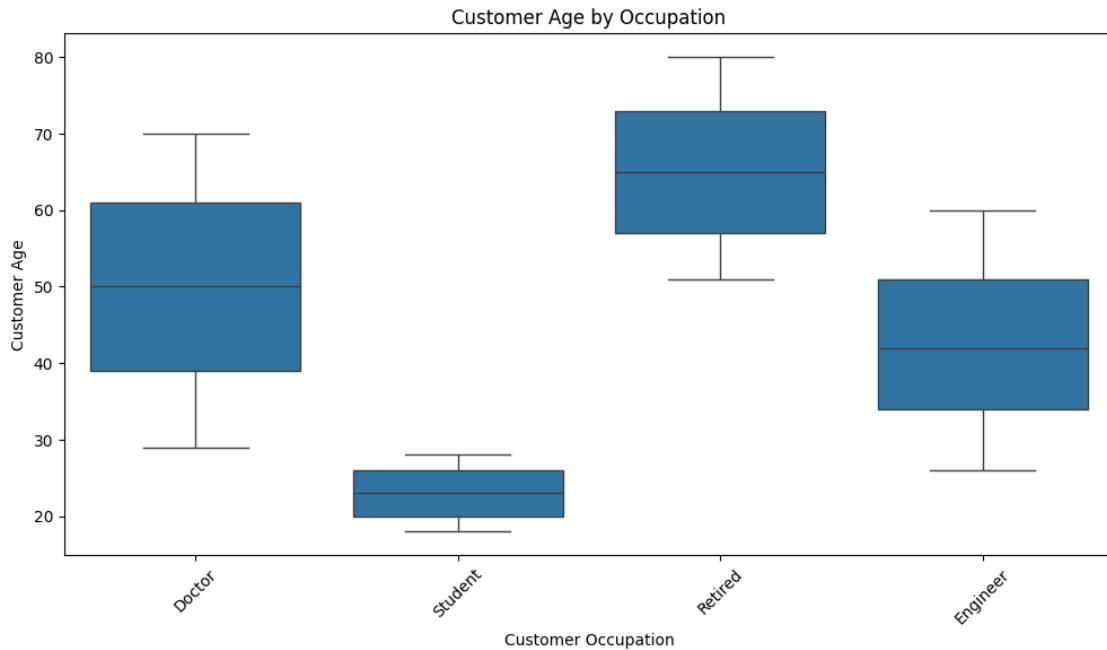
4.7 Login Attempts vs Transaction Amount

```
[10]: # Login Attempts vs Transaction Amount
plt.figure(figsize=(10, 6))
sns.scatterplot(x='LoginAttempts', y='TransactionAmount', hue='TransactionType', data=df)
plt.title('Login Attempts vs Transaction Amount')
plt.xlabel('Login Attempts')
plt.ylabel('Transaction Amount')
plt.legend(title='Transaction Type')
plt.show()
```



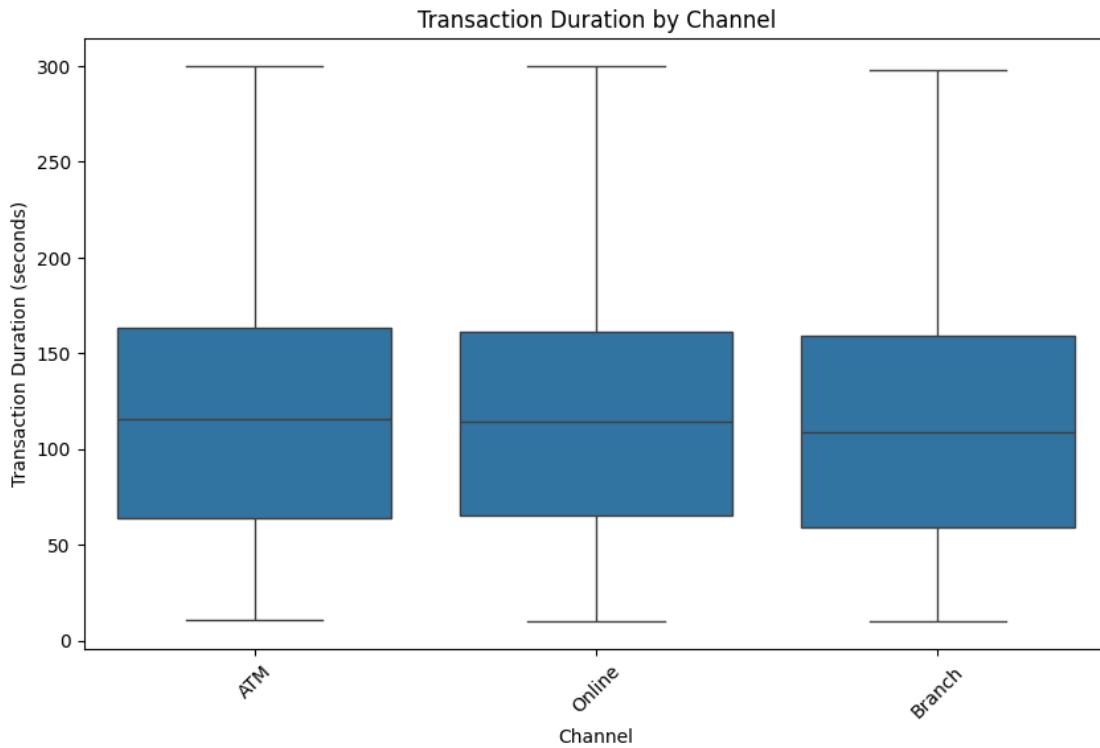
4.8 Customer Age Distribution by Occupation

```
[11]: # Customer Age Distribution by Occupation
plt.figure(figsize=(12, 6))
sns.boxplot(x='CustomerOccupation', y='CustomerAge', data=df)
plt.title('Customer Age by Occupation')
plt.xlabel('Customer Occupation')
plt.ylabel('Customer Age')
plt.xticks(rotation=45)
plt.show()
```



4.9 Transaction Duration by Channel

```
[12]: # Transaction Duration by Channel
plt.figure(figsize=(10, 6))
sns.boxplot(x='Channel', y='TransactionDuration', data=df)
plt.title('Transaction Duration by Channel')
plt.xlabel('Channel')
plt.ylabel('Transaction Duration (seconds)')
plt.xticks(rotation=45)
plt.show()
```



5 Feature Engineering

```
[13]: # Encoding Categorical Variables using Label Encoding
label_encoder = LabelEncoder()
df['TransactionType_Enc'] = label_encoder.fit_transform(df['TransactionType'])
df['Channel_Enc'] = label_encoder.fit_transform(df['Channel'])
df['CustomerOccupation_Enc'] = label_encoder.
    ↪fit_transform(df['CustomerOccupation'])

# Calculating Time Since Last Transaction
df['TimeSinceLastTransaction'] = (df['TransactionDate'] -
    ↪df['PreviousTransactionDate']).dt.total_seconds()
df['TimeSinceLastTransaction'] = df['TimeSinceLastTransaction'].
    ↪fillna(df['TimeSinceLastTransaction'].median())

# Creating Transaction Frequency Feature
transaction_counts = df.groupby('AccountID').size().
    ↪reset_index(name='TransactionCount')
df = pd.merge(df, transaction_counts, on='AccountID', how='left')

# Dropping unnecessary columns for modeling
```

```
df_model = df.drop(['TransactionID', 'AccountID', 'TransactionDate',  

    ↴'TransactionType', 'Location', 'DeviceID',  

        'IP Address', 'MerchantID', 'Channel',  

    ↴'PreviousTransactionDate', 'CustomerOccupation'], axis=1)
```

6 Anomaly Detection using Isolation Forest

```
[14]: # Preparing data for anomaly detection  

features = df_model.columns.drop(['Anomaly']) if 'Anomaly' in df_model.columns  

    ↴else df_model.columns  

X = df_model[features]  

# Initializing and fitting the Isolation Forest model  

isolation_forest = IsolationForest(n_estimators=100, contamination=0.05,  

    ↴random_state=42)  

isolation_forest.fit(X)  

# Predicting anomalies  

df['Anomaly'] = isolation_forest.predict(X)  

df['Anomaly'] = df['Anomaly'].map({1: 0, -1: 1})  

# Number of anomalies detected  

print(f"Number of anomalies detected: {df['Anomaly'].sum()}")  

# Visualizing anomalies  

plt.figure(figsize=(10, 6))  

sns.scatterplot(data=df, x='TransactionAmount', y='AccountBalance',  

    ↴hue='Anomaly', palette={0: 'blue', 1: 'red'})  

plt.title('Anomaly Detection using Isolation Forest')  

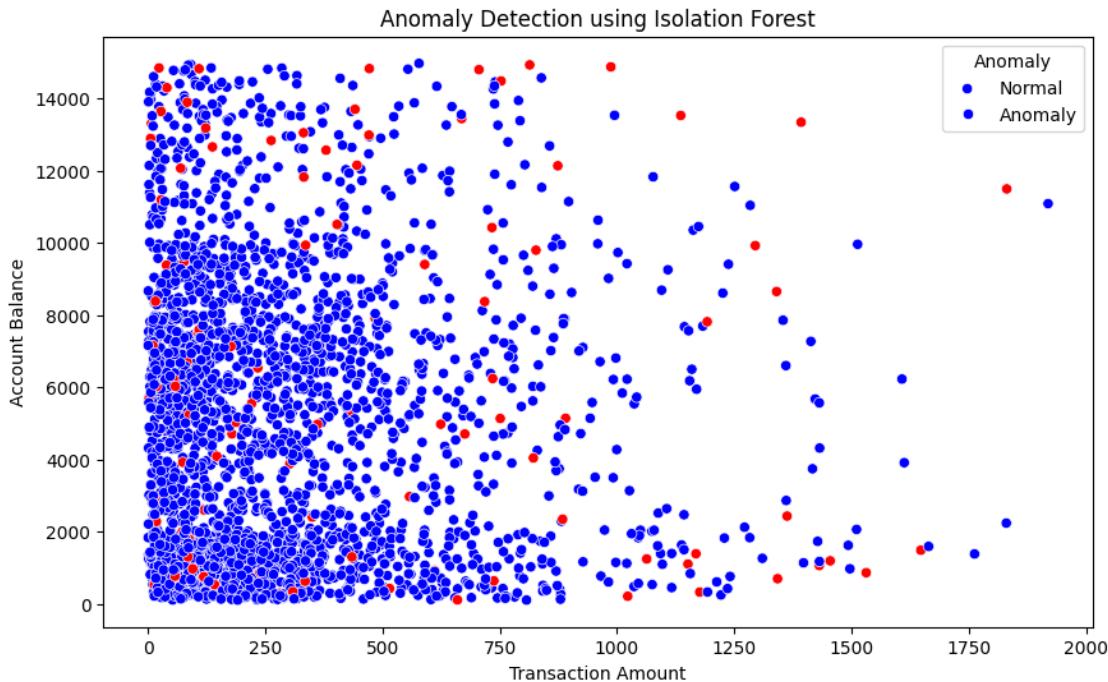
plt.xlabel('Transaction Amount')  

plt.ylabel('Account Balance')  

plt.legend(title='Anomaly', labels=['Normal', 'Anomaly'])  

plt.show()
```

Number of anomalies detected: 126



7 Predictive Modeling

```
[15]: # Using anomalies detected by Isolation Forest as labels
X = df_model[features]
y = df['Anomaly']

# Splitting the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, stratify=y, random_state=42)

# Training a Random Forest Classifier
rf_classifier = RandomForestClassifier(n_estimators=100, random_state=42)
rf_classifier.fit(X_train, y_train)

# Making predictions
y_pred = rf_classifier.predict(X_test)

# Evaluating the model
print("Classification Report:")
print(classification_report(y_test, y_pred))

# Confusion Matrix
conf_matrix = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(6, 4))
```

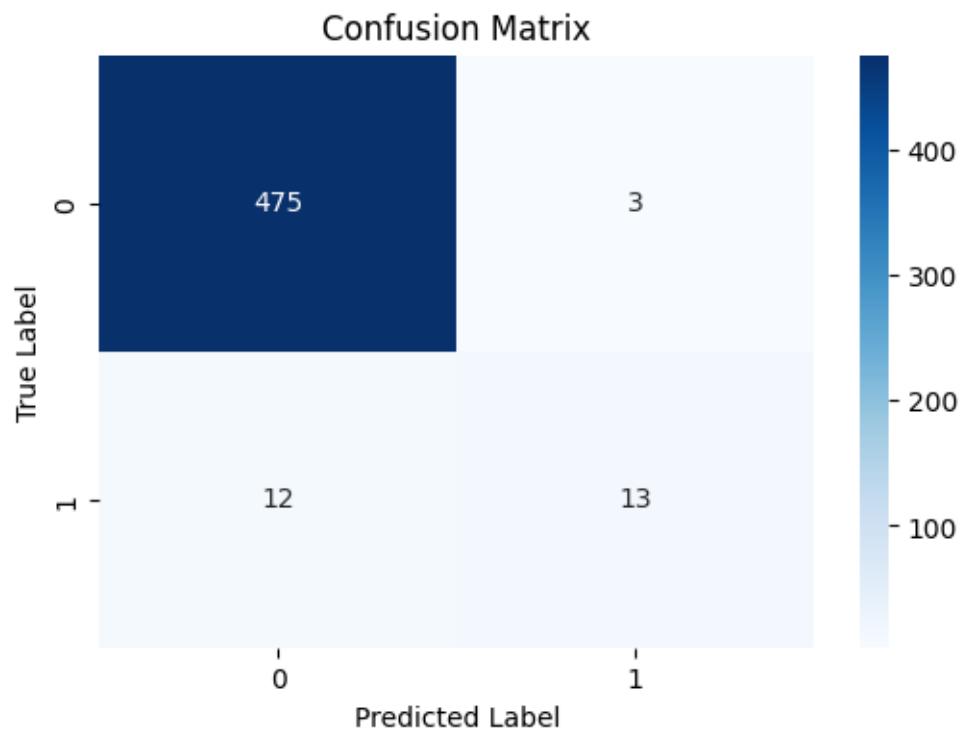
```

sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues')
plt.title('Confusion Matrix')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()

```

Classification Report:

	precision	recall	f1-score	support
0	0.98	0.99	0.98	478
1	0.81	0.52	0.63	25
accuracy			0.97	503
macro avg	0.89	0.76	0.81	503
weighted avg	0.97	0.97	0.97	503



8 Model Improvement and Hyperparameter Tuning

```
[16]: # Hyperparameter tuning using GridSearchCV (optional)
from sklearn.model_selection import GridSearchCV
```

```

# Defining parameter grid
param_grid = {
    'n_estimators': [100, 200],
    'max_depth': [None, 10, 20],
    'max_features': ['auto', 'sqrt']
}

# Initializing GridSearchCV
grid_search = GridSearchCV(estimator=RandomForestClassifier(random_state=42),
                           param_grid=param_grid, cv=3, n_jobs=-1, verbose=2)

# Fitting the model
grid_search.fit(X_train, y_train)

# Best parameters
print(f"Best parameters found: {grid_search.best_params_}")

# Evaluating the best estimator
best_rf = grid_search.best_estimator_
y_pred_best = best_rf.predict(X_test)
print("Classification Report with Best Estimator:")
print(classification_report(y_test, y_pred_best))

```

Fitting 3 folds for each of 12 candidates, totalling 36 fits

```

C:\Users\rakes\AppData\Local\Programs\Python\Python312\Lib\site-
packages\sklearn\model_selection\_validation.py:540: FitFailedWarning:
18 fits failed out of a total of 36.
The score on these train-test partitions for these parameters will be set to
nan.
If these failures are not expected, you can try to debug them by setting
error_score='raise'.

```

Below are more details about the failures:

```

8 fits failed with the following error:
Traceback (most recent call last):
  File "C:\Users\rakes\AppData\Local\Programs\Python\Python312\Lib\site-
  packages\sklearn\model_selection\_validation.py", line 888, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "C:\Users\rakes\AppData\Local\Programs\Python\Python312\Lib\site-
  packages\sklearn\base.py", line 1466, in wrapper
    estimator._validate_params()
  File "C:\Users\rakes\AppData\Local\Programs\Python\Python312\Lib\site-
  packages\sklearn\base.py", line 666, in _validate_params
    validate_parameter_constraints(
  File "C:\Users\rakes\AppData\Local\Programs\Python\Python312\Lib\site-
  packages\sklearn\utils\param_validation.py", line 95, in

```

```

validate_parameter_constraints
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The 'max_features'
parameter of RandomForestClassifier must be an int in the range [1, inf), a
float in the range (0.0, 1.0], a str among {'sqrt', 'log2'} or None. Got 'auto'
instead.

-----
10 fits failed with the following error:
Traceback (most recent call last):
  File "C:\Users\rakes\AppData\Local\Programs\Python\Python312\Lib\site-
packages\sklearn\model_selection\_validation.py", line 888, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "C:\Users\rakes\AppData\Local\Programs\Python\Python312\Lib\site-
packages\sklearn\base.py", line 1466, in wrapper
    estimator._validate_params()
  File "C:\Users\rakes\AppData\Local\Programs\Python\Python312\Lib\site-
packages\sklearn\base.py", line 666, in _validate_params
    validate_parameter_constraints()
  File "C:\Users\rakes\AppData\Local\Programs\Python\Python312\Lib\site-
packages\sklearn\utils\_param_validation.py", line 95, in
validate_parameter_constraints
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The 'max_features'
parameter of RandomForestClassifier must be an int in the range [1, inf), a
float in the range (0.0, 1.0], a str among {'log2', 'sqrt'} or None. Got 'auto'
instead.

    warnings.warn(some_fits_failed_message, FitFailedWarning)
C:\Users\rakes\AppData\Local\Programs\Python\Python312\Lib\site-
packages\sklearn\model_selection\_search.py:1052: UserWarning: One or more of
the test scores are non-finite: [      nan      nan  0.96465877  0.96366449
nan      nan
  0.964162  0.964661      nan      nan  0.96465877  0.96366449]
    warnings.warn(
Best parameters found: {'max_depth': 10, 'max_features': 'sqrt', 'n_estimators':
200}
Classification Report with Best Estimator:
      precision    recall  f1-score   support

          0       0.98      0.99      0.98      478
          1       0.81      0.52      0.63       25

      accuracy                           0.97      503
     macro avg       0.89      0.76      0.81      503
weighted avg       0.97      0.97      0.97      503

```